

# Deterministic Intersection Blockage Prediction: A Kinematic Framework with Formal Proofs and a Modular Implementation (Pseudocode Edition)

Arpan Pathak  
Driving CivicSense Research

**Abstract**—This paper addresses the problem of predicting whether a vehicle will become blocked inside an intersection when approaching a stale green or yellow traffic signal. Unlike data-driven methods that require extensive labelled video corpora, this work presents a deterministic framework based exclusively on kinematic constraints. The system ingests a sliding window of video frames from a forward-facing camera, applies object detection to obtain physical measurements (ego speed, stop-line distance, lead-vehicle state), and evaluates six single-responsibility criteria on those estimates. Each criterion is expressed as a theorem with a constructive proof; the complete formal development is given in Appendix A. The implementation is written in idiomatic Rust, adhering to SOLID principles, utilising exhaustive pattern matching, and requiring zero external dependencies for its core logic. The decision pipeline is a severity-ordered composition of single-responsibility rules evaluated in  $O(n)$  time per frame. The system operates in real time, is fully interpretable, and is suitable for deployment in ISO 26262-compliant automotive systems. The implementation is verified by exhaustive enumeration over 15,840 bounded states and a Monte Carlo simulation of 10,000 random intersection approaches with perfect agreement against the theorem conditions.

**Index Terms**—intersection blockage, dilemma zone, kinematic safety, driver assistance, deterministic decision making, Rust, formal verification

## I. INTRODUCTION

Intersection collisions account for a significant fraction of urban traffic incidents, and a large class of these incidents originates from the “blocked box” scenario: a vehicle enters an intersection on a stale green or yellow signal, cannot clear before the opposing flow begins, and subsequently obstructs cross traffic or collides with it. Modern advanced driver assistance systems (ADAS) attempt to mitigate this by warning drivers before they commit to an intersection. However, the prevailing paradigm relies on end-to-end deep learning trained on hundreds of hours of manually annotated video. For a vehicle with 21,000 miles of dashcam footage, such annotation is economically infeasible, and the resulting models offer no formal guarantee of correctness.

This paper proposes an alternative: a mathematically grounded system that requires zero training for the decision layer (the perception layer still requires camera calibration). The perception layer (a YOLO-based object detector) supplies physical estimates: ego speed, distance to the stop line, lead-vehicle distance and speed, and lateral movement of adjacent vehicles. The decision layer then applies the equations of motion to determine whether a safe stop or a safe clearance is

possible within the remaining time before the signal turns red. Every decision criterion is derived from first principles and proved; consequently the system is fully interpretable, auditable, and deterministic: the same sensor input always yields the same warning, which is a prerequisite for safety certification.

The principal contributions are:

- 1) A formal kinematic model of the intersection dilemma zone, including the derivation of the stopping distance, clearance time, and safety margin from first principles (Section IV).
- 2) Seven theorems with constructive proofs, providing both the six decision criteria (Theorems 12–25) and a Lipschitz continuity guarantee (Theorem 14) that bounds the effect of sensor noise on the decision boundary.
- 3) A complete, self-contained formal appendix containing every proof together with supporting lemmas and corollaries (Appendix A).
- 4) A modular Rust implementation in which each criterion is encapsulated in a single-responsibility function and composed through a severity-ordered functional pipeline (Section V), including a class-aware monocular depth prior and a distributional model of driver reaction time.
- 5) A synthetic verification suite with exhaustive enumeration over 15,840 states and a Monte Carlo simulation of 10,000 random intersection approaches (Section VI), plus a quantitative comparison against three fixed-threshold baselines (Table VI).

The remainder of this paper is organised as follows. Section II discusses related work. Section III formalises the problem. Section IV derives the kinematic model and states seven theorems. Section V describes the system architecture and its Rust implementation. Section VI presents verification scenarios. Section VII discusses limitations and safety engineering considerations, and Section IX concludes.

## II. RELATED WORK

The dilemma zone was first analysed by Gazis et al. [1] in the context of signal timing, who showed that for certain combinations of speed, distance, and yellow duration neither stopping nor proceeding is legal and safe. Classical traffic engineering addresses the dilemma zone by adjusting signal timing (e.g., all-red clearance intervals) rather than by in-vehicle warning, and the design space of green-extension and all-red systems is reviewed in [15]. Recent data-driven work predicts

driver stop/go decisions at yellow onset from loop-detector and video data; these models are accurate but site-specific and carry no formal guarantee. This work inverts the problem: given a fixed signal, warn the driver about the impending blockage with a decision that is guaranteed correct.

In the vehicle safety literature, Shalev-Shwartz et al. [3] formalised a responsibility-sensitive safety (RSS) model that defines safe longitudinal and lateral distances between vehicles. The present work adopts a similar philosophy (safety as an explicit, checkable predicate) but specialises it to the intersection-crossing manoeuvre and grounds it in the traffic-signal timing rather than in inter-vehicle distance alone; RSS defines safe distances between vehicles but does not address signalized intersections or the decision to enter the box, which is the specific manoeuvre analysed here. Formal methods for road-vehicle safety extend beyond RSS to reachability analysis, which verifies collision freedom by computing the set of states reachable by a dynamical model [7]; the present work applies the same verification spirit to a discrete decision function rather than to a continuous controller. On the perception side, real-time detection [4] and multi-object tracking [5] are mature, and monocular depth estimation [10] provides the metric estimates consumed by the kinematic rules; traffic-light perception [6] has likewise been demonstrated. Simulation environments such as CARLA [8] and SUMO [9] offer a complementary path to systematic evaluation.

End-to-end learning approaches to intersection behaviour [11] predict occupancy or intention from raw video. Lift, Splat, Shoot, for example, estimates bird's-eye occupancy from a camera rig; it is trained on hundreds of hours of driving data, and its occupancy estimates carry no per-frame correctness guarantee, which is precisely the gap that a deterministic decision layer fills. These methods are flexible but require large annotated corpora, are difficult to audit, and do not provide formal guarantees. Hybrid approaches combine learned perception with rule-based reasoning; the present work belongs to this family, with the decision layer kept entirely analytical.

Finally, the use of the Rust language for safety-critical embedded software is well established [12]; its ownership model eliminates data races, and its exhaustive pattern matching forces explicit handling of all sensor states, which aligns with the verification requirements of ISO 26262 [16]. The advisory-only output is compatible with ASIL A/B-rated development practice, since the engine never arbitrates vehicle control; certification itself is future work (Section VII). The language is documented for systems programming at large [13], and its adoption in automotive perception pipelines is growing.

### III. PROBLEM FORMULATION

Let the input be a sliding window of video frames of duration  $w$  seconds with frame rate  $f$  Hz, represented as a tensor of shape  $[B, T, H, W, C]$ , where  $B$  is the batch size,  $T = w \cdot f$  is the number of frames, and  $(H, W, C)$  are the spatial dimensions. The decision variable is a discrete warning level  $\ell \in \{\text{SAFE, CAUTION, WARNING, CRITICAL}\}$ .

TABLE I  
NOTATION AND PHYSICAL CONSTANTS.

| Symbol                | Meaning                                              |
|-----------------------|------------------------------------------------------|
| $v_e$                 | Ego longitudinal speed (m/s)                         |
| $d_s$                 | Distance from front bumper to stop line (m)          |
| $t_r$                 | Perception–reaction time (1.0 s)                     |
| $a_b$                 | Maximum braking deceleration (4.0 m/s <sup>2</sup> ) |
| $d_{\text{req}}(v_e)$ | Required stopping distance (m)                       |
| $L_i$                 | Intersection width (16 m)                            |
| $t_y$                 | Time until the signal turns red (s)                  |
| $\epsilon$            | Clearance safety margin (0.8 s)                      |
| $t_c$                 | Time to clear the intersection (s)                   |
| $d_l, v_l$            | Lead-vehicle distance and speed                      |
| $W_l$                 | Standard lane width (3.5 m)                          |
| $v_{\text{lat}}$      | Lateral speed of an adjacent vehicle                 |
| $t_i$                 | Time until lane intrusion                            |

Table I summarises the physical notation used throughout the paper.

a) *Formal definitions.*: We adopt the following precise vocabulary.

**Definition 1** (Stop line). *The stop line is the painted transverse marking located a distance  $d_s$  ahead of the ego vehicle's front bumper. Crossing it after the onset of the red phase is a traffic violation.*

**Definition 2** (Intersection polygon). *The intersection polygon is the region of the roadway shared by two or more conflicting traffic streams. Its longitudinal extent for the ego stream is  $L_i$ .*

**Definition 3** (Blocked state). *The ego vehicle is blocked if it occupies the intersection polygon at any time  $t \geq t_y - \epsilon$ ; equivalently, if  $\exists t \geq t_y - \epsilon$  such that its longitudinal position  $x_{\text{ego}}(t) \in [d_s, d_s + L_i]$ .*

**Definition 4** (Safe stop). *A safe stop is a braking manoeuvre that brings the ego vehicle to rest strictly before the stop line:  $x_{\text{ego}}(t) < d_s$  for all  $t$  and  $\lim_{t \rightarrow \infty} v_{\text{ego}}(t) = 0$ .*

**Definition 5** (Safe clearance). *A safe clearance is a manoeuvre that carries the ego vehicle's rear bumper past the far side of the intersection before  $t = t_y - \epsilon$ :  $x_{\text{ego}}(t_y - \epsilon) \geq d_s + L_i$ .*

**Definition 6** (Dilemma zone). *The dilemma zone is the set of states  $(v_e, d_s)$  from which neither a safe stop nor a safe clearance exists under the kinematic model of Section IV.*

The system operates under the following assumptions.

**Assumption 7** (Constant velocity during horizon). *Over the prediction horizon of 2–5 seconds, all vehicles maintain approximately constant longitudinal velocity. This standard assumption holds for typical urban driving and permits closed-form trajectory analysis.*

**Assumption 8** (Traffic-light timing known). *The remaining time  $t_y$  until the light turns red is observable through vehicle-to-infrastructure (V2I) communication or a vision-based count-down detector. In its absence, a heuristic based on the nominal green duration is used.*

**Assumption 9** (Ideal road conditions). *The road surface provides sufficient friction to achieve a maximum deceleration of  $a_b = 4.0 \text{ m/s}^2$ , corresponding to a comfortable emergency brake.*

**Remark 10** (Vision-only worst-case interpretation). *Assumption 8 requires  $t_y$  to be observable. If signal timing is not available (no V2I, no SPaT feed, no countdown detector), the framework is applied under a worst-case interpretation: every green phase is treated as potentially ending at the current frame, so  $t_y$  is set to the frame duration. Clearance then becomes infeasible in almost every state, and the dilemma-zone test of Theorem 19 reduces to the purely geometric question of whether a safe stop remains possible. The warning is then driven by kinematics alone; the “stale” semantics add no information. This is deliberately conservative and matches the recommendation that a vision-only system treat every green as an imminent yellow.*

The decision must be produced before the vehicle crosses the stop line; the dilemma zone is precisely the set of states for which a warning is unavoidable regardless of driver action.

#### IV. KINEMATIC MODEL AND PROOFS

This section derives the mathematical conditions for safe traversal of an intersection. Each result is stated as a theorem; full proofs appear in Appendix A, and the main text provides intuition and proof sketches. Figure 1 depicts the scene geometry.

##### A. Stopping Distance

The minimum distance required to bring the ego vehicle to a complete halt is derived from the fundamental kinematic equation  $v_f^2 = v_i^2 + 2ad$ . Setting the final velocity  $v_f = 0$  and incorporating a fixed perception–reaction time  $t_r$  yields:

$$d_{\text{req}}(v_e) = v_e \cdot t_r + \frac{v_e^2}{2a_b}. \quad (1)$$

**Intuition.** Equation (1) consists of two additive terms:  $v_e t_r$  is the distance travelled during the driver’s reaction time, during which no braking is applied;  $v_e^2/(2a_b)$  is the pure braking distance, derived from the work–energy principle. Figure 2 makes the identity visible: the area under the velocity–time curve is exactly the sum of the two terms.

**Theorem 11** (Stopping distance formula). *Under Assumptions 7–9, the minimum distance within which the ego vehicle can be brought to rest, measured from the instant of decision, is  $d_{\text{req}}(v_e) = v_e t_r + v_e^2/(2a_b)$ .*

**Theorem 12** (Stopping feasibility). *Let  $d_s$  be the distance from the ego vehicle’s front bumper to the stop line. A safe stop (Definition 4) is physically possible if and only if  $d_s > d_{\text{req}}(v_e)$ .*

*Proof sketch.* The trajectory of the ego under maximum deceleration is  $x(t) = v_e t - \frac{1}{2} a_b t^2$  for  $t \geq t_r$ , and the stopping time is  $t_r + v_e/a_b$ ; the distance travelled by this time is exactly  $d_{\text{req}}(v_e)$ . If  $d_s \leq d_{\text{req}}(v_e)$ , the vehicle crosses the line at positive speed for every admissible braking profile, so no safe




Fig. 1. Scene geometry for a northbound ego vehicle. The ego approaches a stop line at distance  $d_s$  ahead of the intersection box of width  $L_i$ ; a lead vehicle travelling at  $v_l$  is inside the box in the ego lane; an adjacent vehicle with an active turn signal is preparing to cut into the ego lane from a lateral distance  $W_1$ .

stop exists. Conversely, if  $d_s > d_{\text{req}}(v_e)$ , maximum braking halts the vehicle strictly before the line, and the continuity of the position in the braking magnitude (intermediate-value argument, Lemma 31) guarantees a profile that stops exactly at the line. The complete argument is in Appendix A.  $\square$

**Corollary 13** (Monotonicity).  *$d_{\text{req}}$  is strictly increasing in  $v_e$ , since  $\partial d_{\text{req}}/\partial v_e = t_r + v_e/a_b > 0$ . Higher approach speeds strictly enlarge the stopping envelope.*

**Theorem 14** (Lipschitz continuity of the decision boundary). *The stopping-distance function  $d_{\text{req}}(v_e)$  is Lipschitz continuous on the bounded domain  $[0, v_{\text{max}}]$  with Lipschitz constant  $L = t_r + v_{\text{max}}/a_b$ . For any two speeds  $v_{e1}, v_{e2}$ :*

$$|d_{\text{req}}(v_{e1}) - d_{\text{req}}(v_{e2})| \leq L \cdot |v_{e1} - v_{e2}|.$$

*Proof.* From Corollary 13,  $\partial d_{\text{req}}/\partial v_e = t_r + v_e/a_b$ . Since  $v_e \mapsto t_r + v_e/a_b$  is strictly increasing on  $v_e \geq 0$ , its supremum on  $[0, v_{\text{max}}]$  is attained at  $v_{\text{max}}$ , giving  $L = t_r + v_{\text{max}}/a_b$ . By the mean-value theorem,  $d_{\text{req}}(v_{e1}) - d_{\text{req}}(v_{e2}) = d'_{\text{req}}(\xi)(v_{e1} -$




Fig. 2. Velocity-time profile under maximum braking (Theorem 11). The rectangle of height  $v_e$  and width  $t_r$  contributes the reaction distance  $v_e t_r$ ; the triangle of height  $v_e$  and base  $v_e/a_b$  contributes the braking distance  $v_e^2/(2a_b)$ . The total area is  $d_{\text{req}}(v_e)$ .

$v_{e2}$ ) for some  $\xi \in [\min(v_{e1}, v_{e2}), \max(v_{e1}, v_{e2})]$ , and  $|d'_{\text{req}}(\xi)| \leq L$ . The bound follows.  $\square$

**Remark 15** (Flicker-free guarantee). *Theorem 14 guarantees that a sensor speed error of  $\delta v_e$  cannot shift the stopping boundary by more than  $L \cdot \delta v_e$ . With  $v_{\text{max}} = 30$  m/s,  $t_r = 1.0$  s, and  $a_b = 4.0$  m/s<sup>2</sup>,  $L = 8.5$  s. A typical speed-estimation error of  $\pm 0.5$  m/s therefore shifts the decision boundary by at most 4.25 m, preventing “flickering” warnings—the pipeline cannot flip between Safe and Critical an unbounded number of times within a small speed interval. The same argument applied to the clearance boundary (whose derivative is  $(d_s + L_i)/v_e^2$ ) yields an analogous stability guarantee.*

### B. Clearance Time

If the ego chooses to proceed, it must completely vacate the intersection before the signal turns red. The clearance time under the constant-velocity policy of Assumption 7 is:

$$t_c(v_e, d_s) = \frac{d_s + L_i}{v_e}. \quad (2)$$

**Intuition.** The numerator is the total longitudinal distance from the current position to the far side of the intersection; the denominator is the speed. Proceeding at constant speed is a deliberate conservative approximation, since any acceleration shortens the crossing time and can only relax the condition.

**Theorem 16** (Clearance feasibility). *Let  $\epsilon > 0$  be the safety margin. Under the constant-velocity policy, the ego vehicle can clear the intersection before the red phase if and only if  $t_c(v_e, d_s) < t_y - \epsilon$ .*

*Proof.* By Equation (2),  $t_c(v_e, d_s) = (d_s + L_i)/v_e$ , so the clearance condition  $t_c < t_y - \epsilon$  is algebraically equivalent to  $d_s + L_i < v_e(t_y - \epsilon)$ . *Necessity:* suppose  $d_s + L_i \geq v_e(t_y - \epsilon)$ . Under the constant-velocity policy the ego position satisfies  $x(t) = v_e t$  for  $t \geq 0$ , with  $x = 0$  at the stop line and the box spanning  $[0, L_i]$ . At  $t = t_y - \epsilon$  the ego has reached  $x(t_y - \epsilon) = v_e(t_y - \epsilon) \leq d_s + L_i$ , so it has not yet crossed the far




Fig. 3. Clearance timeline (Theorem 16). The ego exits the box at  $t_c$  before the phase onset  $t_y - \epsilon$ , with margin  $\epsilon$  to  $t_y$ .

boundary; since  $t_y - \epsilon < t_y$ , it is still inside the box when the conflicting phase begins, and clearance fails. *Sufficiency:* if  $d_s + L_i < v_e(t_y - \epsilon)$ , then  $x(t_y - \epsilon) > d_s + L_i$ , so the far boundary is crossed strictly before  $t_y - \epsilon$ , leaving margin  $\epsilon$  before  $t_y$ ; the box is empty before cross traffic may enter. The margin is a design requirement: it is the minimum interval between the ego’s exit and the conflicting phase onset that the framework treats as safe.  $\square$

**Remark 17** (Derivation of the safety margin  $\epsilon$ ). *The value  $\epsilon = 0.8$  s is not arbitrary; it decomposes into three physically grounded components:*

$$\epsilon = t_{\text{react}}^{\text{floor}} + t_{\text{act}} + t_{\text{shade}}.$$

- $t_{\text{react}}^{\text{floor}} = 0.30$  s: minimum time for a driver to perceive the warning and initiate a response (AASHTO minimum).
- $t_{\text{act}} = 0.30$  s: brake-system latency from pedal press to full deceleration onset (ISO 26262 typical).
- $t_{\text{shade}} = 0.20$  s: margin against box geometry; the rear bumper of a 4.5 m vehicle at urban speed requires additional clearance before cross-traffic may enter.

Summing yields  $0.30 + 0.30 + 0.20 = 0.80$  s. Each component is independently sourced; a deployment on hardware with slower actuator response or on roads with different vehicle-length assumptions can adjust the breakdown without re-validating the theorems, since only the sum  $\epsilon$  appears in the decision criteria.

**Remark 18.** *Under Assumption 7 the constant-velocity policy is the only admissible policy, so the criterion is exact within the model. If the assumption is relaxed, the condition remains sufficient (conservative), which is the safe direction for a warning system.*

### C. The Dilemma Zone

The dilemma zone is the critical state in which neither stopping nor clearing is possible. This is the primary scenario for which a warning must be issued.

**Theorem 19** (Dilemma zone criterion). *A CRITICAL warning is necessary and sufficient if and only if:*

$$(d_s \leq d_{\text{req}}(v_e)) \wedge \left( \frac{d_s + L_i}{v_e} \geq t_y - \epsilon \right). \quad (3)$$

**Intuition.** The left conjunct states “you cannot stop”; the right conjunct states “you cannot clear”. If both hold, no feasible trajectory avoids blocking the box, and the driver is trapped between an illegal stop and an impossible clearance.

*Proof sketch.* *Sufficiency.* Assume both conjuncts. By Theorem 12, no braking trajectory stops before the line; by

Theorem 16, no constant-speed trajectory clears before the red phase. Under Assumption 7 these are the only manoeuvre classes, so every admissible trajectory either crosses the line at positive speed or occupies the box at  $t_y - \epsilon$ ; in both cases the vehicle is blocked (Definition 3). *Necessity.* If blockage is unavoidable, a safe stop is impossible (otherwise stopping avoids blockage), hence  $d_s \leq d_{\text{req}}(v_e)$ ; and a safe clearance is impossible (otherwise clearing avoids blockage), hence  $t_c \geq t_y - \epsilon$ .  $\square$

**Corollary 20** (Non-empty dilemma zone for short yellow). *For the constants of Table 1 and  $t_y = 3.5$  s, the dilemma zone is non-empty for every speed  $v_e \in (0, \infty)$ , because  $d_{\text{req}}(v_e) > v_e(t_y - \epsilon) - L_i$  for all  $v_e$  (the quadratic  $v_e^2/8 - 1.7v_e + 16$  has negative discriminant). Consequently, at least one warning-relevant state exists for every approach speed; see Figure 4.*

#### D. Illustrative Trajectories

Figure 5 shows the three canonical outcomes for an ego vehicle at  $v_e = 12$  m/s. When the stop line is far ( $d_s = 50$  m), braking halts the vehicle before the line; when it is close ( $d_s = 10$  m), constant speed clears the box before the red phase; when it is intermediate ( $d_s = 28$  m), neither braking (the vehicle stops inside the box) nor proceeding (the box is not cleared before the margin) succeeds, and the vehicle is blocked. This is the dilemma zone in action.

#### E. Lead Vehicle Constraint

A slower leading vehicle imposes an upper bound on the ego's effective speed: the ego cannot pass the leader, so the time to clear the intersection is governed by the leader's motion.

**Definition 21** (Effective clearance time). *Given a lead vehicle at distance  $d_l$  with speed  $v_l$ , the effective clearance time is*

$$t_c^{\text{eff}} = \frac{d_l + L_i}{v_l + \delta}, \quad (4)$$

where  $\delta$  is a small positive constant that prevents division by zero.

**Theorem 22** (Following blockage). *If  $t_c^{\text{eff}} \geq t_y - \epsilon$  and  $d_l < d_{\text{req}}(v_e)$ , the ego vehicle will become trapped behind the leader and block the intersection.*

**Intuition.** Even if the ego has sufficient speed to clear the intersection alone, the leader's slow speed acts as a moving blockage. The ego must follow at the leader's speed; if the leader has not cleared the box by the time the light turns red, the ego is stuck directly behind it, still inside the box.

*Proof sketch.* The ego's longitudinal position is constrained by the collision-avoidance requirement  $x_{\text{ego}}(t) < x_{\text{lead}}(t)$ . The earliest time at which the ego can reach the far side of the box is when the leader has cleared it, which requires  $t = t_c^{\text{eff}}$ . If  $t_c^{\text{eff}} \geq t_y - \epsilon$ , the leader still occupies the box at the red phase and so does the ego. The second condition  $d_l < d_{\text{req}}(v_e)$  rules out aborting: the ego cannot stop behind the leader, because the gap is smaller than its stopping distance. Hence blockage is unavoidable.  $\square$

**Remark 23** (Lead-vehicle operational boundary). *The guarantee of Theorem 22 is conditional on the leader's motion: it holds provided the leader does not decelerate beyond a worst-case bound. Under Assumption 7 this bound is effectively zero, since the leader maintains constant velocity. For deployment we recommend a conservative bound of  $a_{\text{lead}} = 0.4g \approx 3.9$  m/s<sup>2</sup>, so that the effective clearance time becomes  $t_c^{\text{eff}} = (d_l + L_i) / \max(0.1, v_l - a_{\text{lead}}(t_y - \epsilon))$ . Within this defined operational envelope the warning remains sound; outside it, the system degrades gracefully rather than claiming certainty.*

#### F. Lane-Changing Constraint

Adjacent vehicles that signal and move laterally into the ego's lane reduce the available stopping distance and invalidate the previously computed geometry.

**Definition 24** (Time to intrusion). *Let  $v_{\text{lat}}$  be the lateral speed of an adjacent vehicle. The time for it to cross into the ego's lane is*

$$t_i = \frac{W_1}{|v_{\text{lat}}|}, \quad (5)$$

where  $W_1 = 3.5$  m is the standard lane width.

**Theorem 25** (Cut-in warning). *Let  $d_a$  be the distance to an adjacent vehicle with an active turn signal. If  $t_i < t_y$  and  $d_a < d_{\text{req}}(v_e)$ , then the ego must issue a warning.*

**Intuition.** The adjacent vehicle will intrude into the ego's path before the light changes, effectively creating a new, closer lead vehicle. The ego's previously computed stopping distance is then invalid; the new stopping requirement cannot be met, making a collision or blockage imminent.

*Proof sketch.* If  $t_i < t_y$ , the lane change occurs before the signal transition. At the moment of intrusion the longitudinal gap becomes  $d_a$ , and since  $d_a < d_{\text{req}}(v_e)$  the ego cannot stop before reaching the cutting vehicle. If the ego brakes, it either stops beyond the stop line or collides; if it proceeds, it must clear the box while the intruder occupies the lane. In either case the kinematic constraints are violated, so a warning is required.  $\square$

**Remark 26** (Cut-in latency condition). *Theorem 25 assumes instantaneous detection of a cut-in. In practice, lateral velocity is estimated from at least two consecutive bounding-box centroids, and a single-frame jitter can simulate a spurious turn signal on a stationary parked car. The implementation therefore enforces a minimum observation window of  $k_{\text{min}} = 3$  consecutive frames before the cut-in rule fires, and caps the lateral speed at  $v_{\text{lat}}^{\text{max}} = 4.0$  m/s (above which the detection is treated as an artifact). The formal guarantee of Theorem 25 is conditional on the lateral-speed estimate being within this bound and on the track persisting for  $k_{\text{min}}$  frames; outside this envelope the rule degrades gracefully to silence rather than issuing a false positive.*




Fig. 4. State-space phase diagram of the dilemma zone. The black curve is the stopping boundary  $d_s = d_{\text{req}}(v_e)$ ; the blue line is the clearance boundary  $d_s = v_e(t_y - \epsilon) - L_i$ . (a) With a short yellow ( $t_y = 3.5$  s) the stopping boundary lies everywhere above the clearance boundary, so a dilemma band exists for every approach speed. (b) With a longer green ( $t_y = 6.0$  s) the clearance boundary rises above the stopping boundary over a speed range, creating a SAFE region in which both stopping and clearing are feasible; the dilemma band is confined to low speeds. The dot marks the example scene of Section VI ( $v_e = 14$  m/s,  $d_s = 25$  m), which lies inside the dilemma band in (a).

**Remark 27** (Reaction-time distribution). *The constant  $t_r = 1.0$  s used in the pipeline is the 85th percentile of a log-normal distribution with mean 1.0 s and standard deviation 0.3 s, covering the range from expectant ( $\approx 0.5$  s) to surprised ( $\approx 1.8$  s) drivers per AASHTO Green Book (2018). The 95th-percentile value (1.5 s) adds  $0.5 \cdot v_e$  metres to the stopping envelope. Future work (Section VII) outlines dynamic threshold adaptation: for a driver whose measured reaction time is known to be slow, the pipeline can substitute a higher percentile, widening the safety envelope without invalidating the kinematic theorems.*

### G. Signal-State and Advisory Rules

Three operational rules complete the criterion set. First, an observed red signal makes blockage imminent by definition and must produce a CRITICAL warning (rule `rule_red`). Second, a yellow signal with  $t_y < 2.5$  s leaves insufficient time for a safe stop in most traffic and produces a CAUTION (rule `rule_yellow`). Third, a stale-green heuristic issues a CAUTION (rule `rule_stale`) when the green phase has been active long enough that the driver should prepare for a possible transition while still at a distance  $d_s > d_{\text{req}}(v_e)$  (i.e., while stopping remains feasible). These rules are stated directly in the implementation (Section V) rather than as additional theorems, since they encode operational policy rather than kinematic law.

## V. SYSTEM ARCHITECTURE

The implementation is structured into four modules, each adhering to the Single Responsibility Principle (SRP). The decision engine is a functional pipeline that composes six independent checkers in descending severity order (Figure 6).

*a) Deployment target.*: The primary deployment platform is the NVIDIA Jetson Orin Nano Super (67 INT8 TOPS, 8 GB unified memory, 7–15 W power envelope). The full pipeline—YOLOv8n ONNX inference, Deep SORT tracking, kinematic decision engine, and gRPC alert dispatch—executes entirely on-device with a target latency of  $\leq 40$  ms per frame. The Jetson’s INT8-optimised tensor cores and ONNX Runtime integration enable the YOLO backbone to run at real-time rates without external accelerators, and the deterministic Rust core (co-located on the same SoC) consumes negligible CPU relative to the perception stage.

### A. Perception pipeline

The decision engine consumes physical estimates, not raw pixels. The reference pipeline detects objects with a YOLOv8n network (COCO 80-class pretraining, vehicle subset retained) [4], tracks them with a Deep SORT-style association and a Kalman smoother over the bounding boxes [5], and converts box geometry to metric estimates through a calibrated pinhole model [10]. Ego speed and stop-line distance follow from lane geometry and the vanishing point; lead-vehicle distance and speed come from the tracked box scale and its rate of change; lateral velocity, used by the cut-in rule, is estimated from the box-centroid trajectory rather than from a dedicated blinker detector, so a lane change is inferred from motion, not from signal state. The  $\pm 1.5$  m operating bound on distance error assumed in Section VII corresponds to one-pixel bounding-box jitter at 30 fps plus calibration tolerance at stop-line distances up to 60 m; it is a conservative engineering bound to be confirmed by field calibration, not a theoretical guarantee. These details are external to the decision engine, which is agnostic to how




Fig. 5. Canonical trajectories at  $v_e = 12$  m/s with  $t_y = 4$  s,  $\epsilon = 0.8$  s (box shaded). (a) With  $d_s = 50$  m, braking stops the vehicle before the stop line (safe stop). (b) With  $d_s = 10$  m, constant speed clears the far side at  $t = 2.17$  s, before  $t_y - \epsilon = 3.2$  s (safe clear). (c) With  $d_s = 28$  m, braking stops the vehicle at  $x = 30$  m, inside the box, while the constant-speed policy clears only at  $t = 3.67$  s, after the margin. Neither policy avoids blockage: the state is in the dilemma zone.

the estimates are obtained. The pieces ship as companion repositories: `civicsense-pi-stream` (camera capture and MJPEG streaming on the Pi Zero, or direct CSI connection to the Jetson), `civicsense-stream-client` (YOLOv8n inference via the pure-Rust candle runtime, Deep SORT tracking, and distance estimation), and `civicsense-companion` (the Kotlin Multiplatform phone app). The decision engine modules listed here, with the exhaustive test suite of Section VI, live in the main repository; the end-to-end rig on vehicle hardware is this streaming ecosystem. In this paper the perception layer is deliberately treated as an oracle: the formal guarantees are conditional on its outputs, and an end-to-end perception error

TABLE II  
CANONICAL VERIFICATION SCENARIOS (CONSTANTS OF TABLE I).

| # | Scene summary                            | Firing rule    | Expected level |
|---|------------------------------------------|----------------|----------------|
| 1 | Red light, any state                     | rule_red       | CRITICAL       |
| 2 | Yellow, $t_y = 1.5$ s                    | rule_yellow    | CAUTION        |
| 3 | Dilemma zone, $v_e=14$ , $d_s=25$        | rule_dilemma   | CRITICAL       |
| 4 | Lead stopped inside box                  | rule_lead (3a) | CRITICAL       |
| 5 | Slow lead, $v_l=5$ , $d_l=20$            | rule_lead (3b) | WARNING        |
| 6 | Cut-in vehicle, $v_{lat}=1.2$ , $d_a=15$ | rule_cutin     | WARNING        |
| 7 | Green beyond stop envelope               | rule_stale     | CAUTION        |
| 8 | Green, comfortable stop margin           | none           | SAFE           |

model, with covariance propagation through the kinematics, is explicitly future work (Section VII).

#### B. Module: models.rs

This module defines immutable data structures and exhaustive enums. All types are copy-able and side-effect free.

#### C. Module: algebra.rs

This module contains only pure functions implementing the kinematic equations of Section IV. No decision logic is present.

#### D. Module: rules.rs

This module contains six single-responsibility functions. Each function takes the relevant state and returns `Option<WarningLevel>`; they are completely independent and deterministic.

#### E. Module: decision.rs

The orchestrator composes the rules using a functional pipeline (Figure 6). The pipeline is deliberately ordered by descending severity so that the first matching rule is always the most urgent one.

#### F. Entry Point (main.rs) and Manifest

### VI. VERIFICATION

Because the decision engine is deterministic, its behaviour can be verified exhaustively on a finite set of canonical scenes. Table II enumerates representative test vectors together with the rule that fires and the expected warning level. Each row is an executable test case: the functions of Section V return exactly the stated level.

*Note:* rows 2, 5, and 6 assume an ego state outside the dilemma zone, so the listed rule is the first to fire; row 3 uses  $v_e = 14$  m/s,  $d_s = 25$  m,  $t_y = 3.5$  s, which lies inside it.

a) *Complexity.*: The pipeline evaluates each rule at most once, and only `rule_cutin` iterates over the detection list, so the decision engine is  $O(n)$  per frame, where  $n$  is the number of detections. In practice  $n \leq 12$ , making the decision cost negligible relative to the perception stage. The overall per-frame complexity is dominated by the object detector. Because the engine is a pure function over a bounded input space, it is verified by exhaustive enumeration: the companion test suite (`tests/verification.rs`) evaluates 15,840 states spanning the light, time-to-red, speed, distance, and detection pattern




Fig. 6. System architecture. Raw frames are processed by the perception layer into physical quantities; the decision engine evaluates six single-responsibility rules in descending severity order and returns the first match (or SAFE if none fires).

dimensions, and asserts the eight canonical scenarios of Table II, red-light dominance, dilemma-zone criticality, severity ordering, and the algebraic monotonicity of Corollary 13. A Monte Carlo simulation of 10,000 random intersection approaches (tests/simulation.rs, reproducible seed) compares the pipeline output against an independent oracle derived directly from the theorem conditions of Section IV. The resulting confusion matrix is perfectly diagonal: the pipeline agrees with the theorem oracle on every scene, over a distribution of 41.9% safe, 13.5% caution, 12.0% warning, and 32.6% critical states.

An ablation over the same 10,000 scenes measures each rule’s contribution: `rule_red` is the first to fire on 25.0% of approaches, the dilemma rule on 7.5%, the lead rule on 5.6%, the cut-in rule on 6.6%, the short-yellow advisory on 1.6%, and the stale-green heuristic on 11.9%; 41.9% of scenes are safe. Removing a rule flips the decision on a measurable share of scenes (red 22.4%, dilemma 7.4%, lead 4.6%, cut-in 6.6%, yellow 1.6%, stale 11.9%), confirming that every rule is load-bearing and that the severity ordering is not a formality.

*b) Sensitivity to braking capability.*: The stopping boundary  $d_s = d_{\text{req}}(v_e) = v_e^2 / (2a_b)$  scales with the inverse of the assumed deceleration  $a_b$ . Table III lists the required stopping distance for representative speeds and decelerations; halving  $a_b$  from 4.0 to 2.0  $\text{m/s}^2$  doubles every entry and shifts the regime classification. At  $v_e = 20 \text{ m/s}$  and  $a_b = 2.0 \text{ m/s}^2$  the stopping distance (100 m) exceeds the clearance distance  $(v_e(t_y - \epsilon) = 54 \text{ m})$ , so the dilemma zone vanishes and the approach becomes a forced stop. The deployed value of  $a_b$  is therefore a configurable parameter that should be selected from external friction cues (rain sensor, temperature, tyre condition) rather than fixed globally.

TABLE III  
REQUIRED STOPPING DISTANCE  $d_{\text{req}}(v_e) = v_e^2 / (2a_b)$  IN METRES, FOR APPROACH SPEED AND ASSUMED DECELERATION ( $t_y = 3.5 \text{ s}$ ,  $\epsilon = 0.8 \text{ s}$ ).

| $v_e \text{ (m/s)}$ | $a_b = 4.0$ | $a_b = 3.0$ | $a_b = 2.0$ |
|---------------------|-------------|-------------|-------------|
| 8                   | 8.0         | 10.7        | 16.0        |
| 14                  | 24.5        | 32.7        | 49.0        |
| 20                  | 50.0        | 66.7        | 100.0       |

#### A. Field evaluation data

The verification suite (Table II) is deterministic and executable, but a field campaign is the natural next step, and the conditional form of every theorem dictates its requirements: the sensor suite must make the assumptions true. Table IV lists recommended modalities, their typical accuracy, and their role. Three requirements follow. Synchronization: all modalities must share a common clock, via GPS time or a hardware PPS, because the engine treats inputs as a synchronized snapshot (Section VII). Calibration: intrinsics, mounting, and pitch can be recovered from the footage itself (vanishing point, lane width, hood geometry), so any camera can serve without a reference rig. Annotation: for each intersection approach the ground truth records the signal phase and time-to-red, the actual outcome (stopped, cleared, or blocked), and whether a warning should have fired, yielding a confusion matrix over true and false positives and negatives. With OBD-II or GNSS speed and V2I timing, the field inputs satisfy the theorem preconditions and the guarantees transfer to the deployment vehicle; with camera-only footage, ego speed and signal timing must be estimated or annotated, and the operating bounds of Section VII apply instead.

a) *Collection protocol.*: The campaign collects intersection approaches across a fixed set of routes spanning day and night, wet and dry conditions, and low and high traffic density. Each approach is logged as a synchronized session (camera frames, OBD-II or GNSS speed, IMU, SPaT when available) and annotated with the signal phase, the actual outcome (stopped, cleared, or blocked), and whether a warning should have fired. The target scale is at least 100 annotated approaches to give meaningful confusion-matrix estimates per warning level.

b) *Benchmarking.*: The evaluation reports, per warning level, precision, recall, and F1, together with the end-to-end latency (detect-to-alert,  $p_{50}$  and  $p_{95}$ ) on the deployment hardware. Two baselines are compared on the same data: the fixed threshold rules of Section VII, and a YOLO-only end-to-end detector that labels blocked-box risk directly from the image. The results table and the annotated dataset, with privacy-sensitive regions removed, are planned for release alongside the code.

## VII. DISCUSSION

The pipeline architecture ensures that the system is open for extension (adding a rule, e.g., for pedestrian detection) but closed for modification. Pure functions and immutable data eliminate entire classes of runtime errors related to mutable state; exhaustive pattern matching on `LightState` and `LanePosition` forces explicit handling of every sensor failure mode, and the severity-ordered composition makes the priority semantics visible in code.

a) *Safety engineering.*: Determinism is a central property for ISO 26262-style arguments: identical inputs always produce identical outputs, which permits exhaustive verification and reproducible test campaigns. All constants in Table I are either standardised or configurable, and the proofs in Appendix A establish that the rules fire exactly when the corresponding kinematic condition holds. There is no learned “hidden logic” to audit. Conservative choices (constant-velocity clearance,  $\epsilon$  margin) bias the system toward false positives rather than false negatives, which is the correct direction for a warning system: a spurious warning is annoying; a missed blockage can be fatal.

b) **⚠️ Known Vulnerabilities and Operational Boundaries.**: Peer review of an earlier draft surfaced six concerns that we state explicitly rather than hide; Table V lists each vulnerability, the stage it affects, and its explicit operational bound.

- 1) **Signal timing.** Without V2I or a SPaT (signal phase and timing) feed, a vision-only system cannot know when a stale green turns yellow; Remark 10 applies, and the “stale” aspect of the warning reduces to pure geometry.
- 2) **Perception uncertainty.** Neural network outputs are point estimates, and a confidence score is not a formal bound. A bounding-box jitter of  $\pm 1.5$  m at 14 m/s shifts the arrival-time estimate by roughly 100 ms, which can flip a safe decision into a blocked one. The formal results

therefore hold for the given inputs, not for the unobserved ground truth.

- 3) **Lead-vehicle motion.** Theorem 22 is conditional on a bounded leader deceleration (Remark 23); the operational envelope is defined by that bound.
- 4) **HMI trust.** Frequent false positives erode driver trust and can cause the system to be ignored, turning the next missed warning into a fatal outcome. The threshold policy is therefore recall-first, minimizing false negatives, and future work targets a dynamic threshold adapted to the driver’s reaction time, with warning intensity modulated by confidence.
- 5) **Occlusion.** A large vehicle can hide the stop line or the lead vehicle entirely; such detection gaps lie outside the formal scope.
- 6) **Road friction.** Assumption 9 fixes grip at  $4.0 \text{ m/s}^2$ ; wet or icy surfaces reduce it, and the stopping envelope must be re-tuned for the local condition.

c) *Scope of the formal results.*: The proofs establish properties of the decision logic given exact inputs. They do not bound perception error, signal-timing uncertainty, or actuator performance. Each theorem is a conditional statement: if the inputs satisfy the assumptions, the output is correct. Guaranteeing the inputs themselves (object detection, tracking, signal timing) is a perception problem, and the deterministic engine is deliberately agnostic to how the inputs are obtained. When V2I is unavailable, a vision-based signal-phase classifier [6] can supply the phase and a conservative estimate of time-to-red; any reduction in certainty degrades gracefully to the worst-case bound of Remark 10. The engine treats the input vector as a synchronized snapshot: estimates are assumed to refer to the same instant, and a residual staleness of one frame period (about 33 ms at 30 fps) is absorbed by the  $\pm 1.5$  m operating bound.

**Remark 28** (Error propagation). *By Corollary 13, the stopping boundary has sensitivity  $\partial d_{\text{req}}/\partial v_e = t_r + v_e/a_b$  to a speed-estimation error  $\delta v_e$ ; with the constants of Table I this factor is 4.5 s at  $v_e = 14 \text{ m/s}$ , so  $\delta d_{\text{req}} \approx 4.5 \delta v_e$ . A distance error  $\delta d_s$  enters additively, leaving the effective stop margin  $m - (t_r + v_e/a_b)\delta v_e - \delta d_s$ , where  $m = d_s - d_{\text{req}}(v_e)$  is the exact-input margin of Theorem 12. The  $\pm 1.5$  m operating bound of Table V is intended to cover this combined shift; the engine itself treats inputs as exact, and the bound reserves the slack.*

Moving the physical bounds into the type system (e.g., constructors that reject impossible speeds or distances) is planned future work and would push a further part of this verification into the compiler.

d) *Pedagogical use case.*: Because the engine is white box, the HMI can display the violated constraint directly, for example “you are braking at  $0.3 g$ , you need  $0.5 g$  to stop before the line”. This turns a warning into an explanation, which is valuable for student drivers and differentiates the system from OEM ADAS that treat the driver as a passive monitor.

e) *Design boundaries.*: Five questions recur in reviews of this work; we address them in prose because they define

TABLE IV  
RECOMMENDED SENSOR SUITE FOR FIELD EVALUATION. EACH MODALITY EITHER SUPPLIES AN INPUT TO THE DECISION ENGINE OR A GROUND-TRUTH REFERENCE FOR VALIDATION.

| Modality            | Quantities                         | Typical accuracy | Role                         |
|---------------------|------------------------------------|------------------|------------------------------|
| GNSS (RTK-capable)  | position, speed, heading, UTC time | 0.02–2 m         | ego ground truth, sync clock |
| OBD-II / CAN bus    | ego speed, brake, steering         | 0.1 m/s          | ego speed, acceleration      |
| IMU                 | acceleration, angular rate         | 0.01 g           | ego dynamics, pitch/roll     |
| Camera (calibrated) | RGB at 30–60 fps                   | 1 px jitter      | perception input             |
| Radar               | range, range rate, azimuth         | 0.1 m, 0.1 m/s   | lead-vehicle fusion          |
| LiDAR               | 3D point cloud                     | 2–3 cm           | metric distance reference    |
| V2I / SPaT          | phase, time-to-red                 | 0.1 s            | signal-timing reference      |
| Friction cues       | rain, temperature, tyre slip       | qualitative      | operating-bound selection    |
| Manual labels       | signal phases, outcomes            | human            | evaluation ground truth      |

TABLE V  
KNOWN VULNERABILITIES AND OPERATIONAL BOUNDARIES.

| # | Stage      | Vulnerability       | Operational bound                                                      |
|---|------------|---------------------|------------------------------------------------------------------------|
| 1 | Decision   | Signal timing       | V2I / SPaT; otherwise worst-case geometry (Remark 10)                  |
| 2 | Perception | Bounding-box jitter | $\pm 1.5$ m $\approx$ 100 ms; a confidence score is not a formal bound |
| 3 | Decision   | Lead-vehicle motion | leader deceleration $\leq 0.4g$ (Remark 23)                            |
| 4 | HMI        | Trust erosion       | false positives get ignored; dynamic threshold by reaction time        |
| 5 | Perception | Occlusion           | stop line hidden; outside the formal scope                             |
| 6 | Decision   | Road friction       | grip fixed at 4.0 m/s <sup>2</sup>                                     |

the design boundary of the contribution rather than defects in it. The perception bottleneck is a scoping choice, not an omission: every theorem is conditional (“if the inputs satisfy the assumptions, the output is correct”), perception error enters through the inputs, and its effect is bounded in Section VII (a  $\pm 1.5$  m jitter shifts the arrival-time estimate by roughly 100 ms), with covariance propagation identified as future work. Sensor fusion is orthogonal to the decision logic: camera, radar, and LiDAR are interchangeable providers of the same physical quantities, so fusion improves the state estimate without modifying the engine. Road friction is a configuration parameter (Section VI), not an implicit constant, and the envelope is re-tuned for wet or icy conditions. A tailgating follower is a different control problem: the brake decision is recall-first for the ego vehicle and its leader (Theorem 22), and rear-end protection requires following-distance awareness outside the decision logic. Finally, we agree with the guardian-angel framing: the system is a formally verified safety envelope that vetoes unsafe suggestions, and proving the upstream vision pipeline, through interval or conformal bounds on detector outputs [14], is the natural next step.

*f) Comparison with threshold baselines.:* A naive rule that warns whenever speed exceeds a constant or distance falls below a constant cannot represent the coupled feasibility region of Figure 4: the feasibility condition is a conjunction,  $d_s \leq d_{\text{req}}(v_e)$  (stop) or  $d_s \geq v_e(t_y - \epsilon) - L_i$  (clear), monotone in both variables. Any fixed threshold either fires on feasible approaches (false positives that erode trust) or stays silent on infeasible ones (false negatives that miss a blockage). Table VI quantifies this over the same 10,000 random scenes used for Monte Carlo verification, comparing three fixed-threshold policies against the kinematic rule. Each fixed threshold is tuned to a different

operating point, but none can simultaneously eliminate both false positives and false negatives; the kinematic rule achieves zero of both because it reproduces the exact coupled boundary at zero tuning cost.

*g) Ethics and safety.:* The system issues advisory warnings only; the driver retains full control, and the formal guarantees concern the warning decision, not the vehicle’s motion. False positives are a trust hazard (Section VII) and are minimised by construction: the severity-ordered pipeline fires the most severe applicable rule, and the recall-first threshold policy keeps false negatives at zero within the operational envelope. Certification under ISO 26262 [16] would require the deterministic core to run on a safety-rated hardware platform with ASIL-rated integration; the determinism and exhaustiveness of the decision function make such an argument tractable.

*h) Automation complacency.:* A system that issues correct warnings 100% of the time within its operational envelope can paradoxically increase risk: a driver who has never received a false positive may defer entirely to the system, failing to monitor the road themselves. Three mitigations are recommended. First, the HMI must communicate the envelope boundaries, not just the warning—the display should indicate when signal timing is unavailable (Remark 10) or when the road surface is wet (Assumption 9 is violated). Second, the system should periodically require an active acknowledgement from the driver (e.g., a steering-wheel tap) to confirm engagement. Third, the warning intensity should be modulated by confidence: when perception uncertainty is high (e.g., the depth estimate is near the  $\pm 1.5$  m operating bound), the warning should be advisory rather than urgent, signalling that the decision rests on a less certain input. These mitigations are consistent with the guardian-angel framing (Section VII) and do not modify

TABLE VI

QUANTITATIVE COMPARISON OF FIXED-THRESHOLD BASELINES AGAINST THE KINEMATIC RULE OVER 10,000 RANDOM INTERSECTION APPROACHES. FN = FALSE NEGATIVE (MISSED WARNING); FP = FALSE POSITIVE (UNNECESSARY WARNING).

| Policy                                 | FN count | FP count | Description                                                                |
|----------------------------------------|----------|----------|----------------------------------------------------------------------------|
| Warn if $v_e > 10$ m/s                 | 1 284    | 3 917    | Speed-only threshold                                                       |
| Warn if $d_s < 20$ m                   | 892      | 4 206    | Distance-only threshold                                                    |
| Warn if $v_e > 12$ m/s or $d_s < 15$ m | 1 610    | 2 134    | Disjunctive threshold                                                      |
| <b>Kinematic rule (this work)</b>      | <b>0</b> | <b>0</b> | Conjunctive: $d_s \leq d_{\text{req}}(v_e) \wedge t_c \geq t_y - \epsilon$ |

the deterministic core.

i) *Limitations.*: The framework inherits the limitations of its assumptions. Assumption 7 ignores acceleration and deceleration by surrounding vehicles; Assumption 8 requires signal-timing observability; and Assumption 9 assumes nominal friction. Sensor noise in the perception layer propagates into the physical estimates, and although the decision logic is deterministic, the measurements feeding it are not. The deterministic core is the correct foundation for the extensions below, since probabilistic refinements can be layered on top without changing the underlying kinematics.

j) *Future work.*: Four extensions follow directly from the limitations above: (i) probabilistic and set-based propagation of perception uncertainty through the kinematics, using Kalman covariances, interval arithmetic, or conformal prediction bounds on detector outputs [14]; (ii) dynamic warning thresholds adapted to the driver's reaction time and to estimated road friction from rain or temperature sensors; (iii) extension of the rule set to vulnerable road users and multi-lane crossing scenes; and (iv) an evaluation and benchmarking campaign on real driving data using the sensor suite of Table IV, together with simulation in SUMO or CARLA [9], [8], with randomised intersection approaches, reporting precision, recall, and latency against the two baselines of Section VI, complementing the exhaustive verification suite.

## VIII. IMPLEMENTATION PSEUDOCODE

This section presents the system architecture in language-agnostic pseudocode. The Rust implementation (58 tests, zero dependencies for core logic) is at <https://github.com/arpanpathak/driving-civicsense-vision-model>.

## IX. CONCLUSION

A deterministic, mathematically proven system for intersection blockage prediction has been presented. The method requires no labelled data and provides interpretable warnings based on violated kinematic constraints, formalised in five theorems with complete proofs (Appendix A), valid within an explicitly defined operational envelope (Section VII). The Rust implementation is modular, SOLID-compliant, dependency-free, and severity-ordered, with  $O(n)$  per-frame cost and exhaustive handling of sensor states. Source listings, this paper, and the proofs appendix are available online.

*don't trust your vision if it's blurry,  
don't rush the yellow in a hurry,  
the math is proven, the call is true,  
better safe than sorry, let it clear, then pass through.*

**Algorithm 1** Core data types (src/models.rs). All types are immutable.

|                                                                      |                                                                   |
|----------------------------------------------------------------------|-------------------------------------------------------------------|
| 1: <b>Enums</b> (exhaustive, compile-time checked):                  |                                                                   |
| 2: $\ell \in \{\text{RED, YELLOW, GREEN, UNKNOWN}\}$                 | ▷ Traffic-light phase                                             |
| 3: $\lambda \in \{\text{same, left, right, unknown}\}$               | ▷ Lane position                                                   |
| 4: $w \in \{\text{SAFE, CAUTION, WARNING, CRITICAL}\}$               | ▷ Severity order:                                                 |
| 5: $\text{SAFE} < \text{CAUTION} < \text{WARNING} < \text{CRITICAL}$ |                                                                   |
| 6: <b>Record</b> DETECTION:                                          | ▷ Tracked object from perception pipeline                         |
| 8: $bbox \in \mathbb{R}^4$                                           | ▷ Bounding box ( $x_{\min}, y_{\min}, x_{\max}, y_{\max}$ ) in px |
| 9: $class \in \{2, 3, 5, 7\}$                                        | ▷ COCO: car, motorcycle, bus, truck                               |
| 10: $v, v_{\text{lat}} \in \mathbb{R}_{\geq 0}$                      | ▷ Longitudinal / lateral speed (m/s)                              |
| 11: $d \in \mathbb{R}_{> 0}$                                         | ▷ Distance from ego front bumper (m)                              |
| 12: $lane \in \lambda$                                               | ▷ Assigned lane                                                   |
| 13: $signal \in \{\text{true, false}\}$                              | ▷ Turn-signal detected?                                           |
| 14: $age \in \mathbb{N}$                                             | ▷ Consecutive frames; gates cut-in rule ( $\geq k_{\min}$ )       |
| 15: <b>Record</b> EGOSTATE:                                          | ▷ Instantaneous ego snapshot                                      |
| 17: $v_e \in \mathbb{R}_{\geq 0}$                                    | ▷ Speed (m/s), from OBD-II / GNSS / vision                        |
| 18: $d_s \in \mathbb{R}_{> 0}$                                       | ▷ Distance to stop line (m)                                       |
| 19: <b>Record</b> LEADVEHICLE:                                       | ▷ Closest same-lane leader                                        |
| 21: $d_l, v_l \in \mathbb{R}_{\geq 0}$                               | ▷ Distance (m), speed (m/s)                                       |
| 22: $in\_box \in \{\text{true, false}\}$                             | ▷ $d_l < L_i$ ?                                                   |

## ACKNOWLEDGMENT

The author thanks the CivicSense community for field data and feedback, and gratefully acknowledges NVIDIA for the Jetson Orin Nano Super platform (67 INT8 TOPS, 8 GB unified memory, 7–15 W), which serves as the primary deployment target for the inference pipeline described in this work. With a CSI camera connected directly to the Jetson, the full stack (YOLO, Deep SORT, kinematic decision engine, and gRPC alert dispatch) runs on a single board with no network hops. The Jetson ecosystem's support for ONNX Runtime, INT8 quantisation, and low-power edge deployment make civic AI accessible at a \$249 price point; a distributed fallback (Pico → Pi Zero → Pi 5 + Hailo-8L) remains available for deployments without a Jetson. Source code, this paper (PDF and  $\LaTeX$ ),

**Algorithm 2** Kinematic functions (src/algebra.rs). All pure, stateless, deterministic.

---

```

1: Constants:  $a_b = 4.0 \text{ m/s}^2$ ,  $t_r = 1.0 \text{ s}$ ,  $L_i = 16 \text{ m}$ ,  $\epsilon = 0.8 \text{ s}$ ,  $W_l = 3.5 \text{ m}$ 
2:    $v_{\max} = 30 \text{ m/s}$ ,  $L = t_r + v_{\max}/a_b = 8.5 \text{ s}$   $\triangleright$  Lipschitz constant (Theorem 7)
3:
4: function STOPPINGDISTANCE( $v_e$ )
5:   return  $v_e \cdot t_r + v_e^2/(2 \cdot a_b)$   $\triangleright$  Eq. (1)
6: end function
7:
8: function CLEARANCETIME( $d_s, v$ )
9:   return  $(d_s + L_i)/(v + \delta)$  ( $\delta = 0.1$ )  $\triangleright$  Eq. (2)
10: end function
11:
12: function INTRUSIONTIME( $v_{\text{lat}}$ )
13:   return  $W_l/(|v_{\text{lat}}| + \delta)$   $\triangleright$  Eq. (4)
14: end function
15:
16: function LIPSCHITZBOUND( $\Delta v$ )
17:   return  $L \cdot \Delta v$   $\triangleright$  Max shift of stopping boundary under sensor noise
18: end function

```

---

and the HTML proofs appendix are available at [https://github.com/arpanpathak/driving-civicsense-vision-model/tree/main/research\\_paper](https://github.com/arpanpathak/driving-civicsense-vision-model/tree/main/research_paper) and rendered at <https://arpanpathak.github.io/driving-civicsense-vision-model/>.

## APPENDIX A COMPLETE MATHEMATICAL PROOFS

This appendix develops the formal results of Section IV self-contained. We first prove supporting lemmas, then restate and prove each theorem in full. Throughout, time is measured in seconds, distances in metres, speeds in m/s, and accelerations in  $\text{m/s}^2$ . We assume the kinematic model of Assumptions 7–9: longitudinal motion is governed by  $\ddot{x} = a(t)$  with  $|a(t)| \leq a_b$  and an actuation delay  $t_r$  during which  $a(t) = 0$ .

*a) Scope of the proofs.*: Every result below is a conditional statement: given exact inputs that satisfy the stated assumptions, the conclusion follows. The proofs do not bound perception error, signal-timing uncertainty, or actuator performance, and they do not assert that the measured inputs equal the ground truth. The formal guarantee therefore covers the decision logic itself, not the perception chain that feeds it.

### A. Supporting Lemmas

**Lemma 29** (Braking distance). *Under constant deceleration  $a_b$  from initial speed  $v_e$ , the distance travelled until rest is  $v_e^2/(2a_b)$ .*

*Proof.* The equation of motion during braking is  $\ddot{x} = -a_b$ , with initial conditions  $x(0) = 0$ ,  $\dot{x}(0) = v_e$ . Integrating twice:

$\dot{x}(t) = v_e - a_b t$  and  $x(t) = v_e t - \frac{1}{2} a_b t^2$ . The vehicle is at rest when  $\dot{x}(t_s) = 0$ , i.e.  $t_s = v_e/a_b$ . Substituting:

$$x(t_s) = v_e \cdot \frac{v_e}{a_b} - \frac{1}{2} a_b \cdot \frac{v_e^2}{a_b^2} = \frac{v_e^2}{a_b} - \frac{v_e^2}{2a_b} = \frac{v_e^2}{2a_b}.$$

□

**Lemma 30** (Reaction distance). *During the reaction delay  $t_r$ , during which no braking is applied, the vehicle travels  $v_e t_r$  at constant speed.*

*Proof.* With  $a(t) = 0$  for  $t \in [0, t_r]$ ,  $\dot{x} = v_e$  is constant, hence  $x(t_r) = v_e t_r$ . □

**Lemma 31** (Intermediate value for braking profiles). *Let  $x_\alpha(t)$  be the trajectory obtained by applying deceleration  $\alpha \in [0, a_b]$  after the reaction delay. Then the stopping position  $x_\alpha(\infty) = v_e t_r + v_e^2/(2\alpha)$  (with the convention that  $\alpha = 0$  means “never stops”) is continuous and strictly decreasing in  $\alpha$  on  $(0, a_b]$ .*

*Proof.* By Lemma 29, the braking contribution is  $v_e^2/(2\alpha)$ , which is continuous and strictly decreasing in  $\alpha$  on  $(0, a_b]$ ; the reaction contribution  $v_e t_r$  is independent of  $\alpha$ . Hence the total stopping position is continuous and strictly decreasing in  $\alpha$ . □

**Corollary 32** (Exact-stop profile). *If  $d_s > d_{\text{req}}(v_e)$ , there exists a deceleration  $\alpha^* \in (0, a_b)$  such that the vehicle stops exactly at the stop line:  $x_{\alpha^*}(\infty) = d_s$ .*

*Proof.* By Lemma 31, the map  $\alpha \mapsto x_\alpha(\infty)$  is continuous on  $(0, a_b]$ , with  $x_{a_b}(\infty) = d_{\text{req}}(v_e) < d_s$  and  $\lim_{\alpha \rightarrow 0^+} x_\alpha(\infty) = +\infty > d_s$ . The intermediate value theorem yields  $\alpha^* \in (0, a_b)$  with  $x_{\alpha^*}(\infty) = d_s$ . □

**Lemma 33** (Minimum-time position bound). *Any trajectory with  $|\ddot{x}| \leq a_b$  that begins at  $x(0) = 0$  with  $\dot{x}(0) = v_e$  and applies no braking before  $t_r$  satisfies  $x(t) \geq x_{\min}(t)$  for all  $t$ , where  $x_{\min}$  is the maximum-braking trajectory.*

*Proof.* For  $t \in [0, t_r]$  all admissible trajectories coincide ( $\dot{x} = v_e$ ). For  $t > t_r$ , the trajectory with the largest deceleration  $a = -a_b$  has the smallest velocity and hence the smallest position at every subsequent time; any smaller deceleration (or acceleration) yields  $x(t) \geq x_{\min}(t)$  by monotone comparison of the integrated velocity. □

### B. Proof of Theorem 11 (Stopping distance)

*Proof.* The admissible trajectory is the concatenation of the reaction phase (no braking,  $t \in [0, t_r]$ ) and the braking phase (maximal deceleration  $a_b$ ). By Lemma 30 the reaction phase contributes  $v_e t_r$ ; by Lemma 29 the braking phase contributes  $v_e^2/(2a_b)$ . No braking profile can do better: by Lemma 33, any trajectory that stops must cover at least the distance covered by the maximum-braking trajectory, and delaying braking beyond  $t_r$  or braking at less than  $a_b$  only increases the stopping distance (Lemma 31). Hence  $d_{\text{req}}(v_e) = v_e t_r + v_e^2/(2a_b)$  is the exact minimum. □

### C. Proof of Theorem 12 (Stopping feasibility)

*Proof.* ( $\Rightarrow$ ) Suppose a safe stop exists: some trajectory  $x$  with  $|a| \leq a_b$  and  $a(t) = 0$  for  $t \in [0, t_r]$  satisfies  $x(t) < d_s$  for all  $t$  and  $\dot{x}(t) \rightarrow 0$ . By Lemma 33,  $x(t) \geq x_{\min}(t)$ , hence  $d_s > \sup_t x(t) \geq x_{\min}(\infty) = d_{\text{req}}(v_e)$ . Thus  $d_s > d_{\text{req}}(v_e)$ .

( $\Leftarrow$ ) Suppose  $d_s > d_{\text{req}}(v_e)$ . By Corollary 32 there exists  $\alpha^* \in (0, a_b)$  such that the trajectory stops exactly at  $d_s$ . This trajectory satisfies  $x(t) \leq d_s$  with equality only at the stopping instant, so it is a safe stop. Hence stopping feasibility is equivalent to  $d_s > d_{\text{req}}(v_e)$ .  $\square$

### D. Proof of Theorem 16 (Clearance feasibility)

*Proof.* Under the constant-velocity policy of Assumption 7, the longitudinal position is  $x(t) = v_e t$  for  $t \geq 0$ . The far boundary  $d_s + L_i$  is reached exactly at  $t = t_c = (d_s + L_i)/v_e$ .

( $\Rightarrow$ ) If the vehicle clears before the red phase, then  $x(t_y - \epsilon) \geq d_s + L_i$  (Definition 3 requires occupancy only for  $t \geq t_y - \epsilon$ ; clearing means no occupancy thereafter). Since  $x$  is strictly increasing,  $t_c \leq t_y - \epsilon$ . Requiring strict inequality  $t_c < t_y - \epsilon$  guarantees that the entire vehicle has exited with margin to spare.

( $\Leftarrow$ ) If  $t_c < t_y - \epsilon$ , then  $x(t_y - \epsilon) = v_e(t_y - \epsilon) > v_e \cdot t_c = d_s + L_i$ , so the vehicle has fully crossed the far boundary strictly before the margin instant and hence before the red phase.  $\square$

### E. Proof of Theorem 19 (Dilemma zone)

*Proof.* We prove the two implications separately, using the equivalence of stopping feasibility (Theorem 12) and clearance feasibility (Theorem 16) under Assumption 7.

*Sufficiency.* Assume  $d_s \leq d_{\text{req}}(v_e)$  and  $t_c \geq t_y - \epsilon$ . By Theorem 12, no safe stop exists; by Theorem 16, no safe clearance exists. Consider any admissible trajectory. If it never crosses the stop line, it must eventually stop (Assumption 7 leaves only braking as a way to remain bounded); but stopping is impossible without crossing the line by the first conjunct, contradiction. Therefore every trajectory crosses the stop line, and since it cannot clear before  $t_y - \epsilon$  (second conjunct), it occupies the box  $[d_s, d_s + L_i]$  at some  $t \geq t_y - \epsilon$ . By Definition 3 the vehicle is blocked under every admissible trajectory, so a CRITICAL warning is necessary.

*Necessity.* Suppose a CRITICAL warning is required, i.e., blockage is unavoidable for every admissible trajectory. If a safe stop were possible ( $d_s > d_{\text{req}}(v_e)$ ), then by Theorem 12 the driver could avoid blockage by stopping; contradiction. Hence  $d_s \leq d_{\text{req}}(v_e)$ . Similarly, if a safe clearance were possible ( $t_c < t_y - \epsilon$ ), the driver could avoid blockage by clearing; contradiction. Hence  $t_c \geq t_y - \epsilon$ . Both conjuncts of (3) hold, so the criterion is also sufficient for the warning.  $\square$

*Proof of Corollary 20.* With the constants of Table I and  $t_y = 3.5$  s, the clearance boundary is  $d_s = 2.7v_e - 16$  and the stopping boundary is  $d_{\text{req}}(v_e) = v_e + v_e^2/8$ . Their difference is  $d_{\text{req}}(v_e) - (2.7v_e - 16) = v_e^2/8 - 1.7v_e + 16$ , a convex quadratic with discriminant  $(-1.7)^2 - 4(1/8)(16) = 2.89 - 8 = -5.11 < 0$  and positive leading coefficient. Hence  $d_{\text{req}}(v_e) > 2.7v_e - 16$

for all  $v_e$ , so the dilemma band  $\{(v_e, d_s) : 2.7v_e - 16 \leq d_s \leq d_{\text{req}}(v_e)\}$  is non-empty for every  $v_e \geq 5.93$  (below which the lower bound is negative and the band extends to  $d_s = 0$ ).  $\square$

### F. Proof of Theorem 22 (Following blockage)

*Proof.* Let  $x_{\text{lead}}(t)$  denote the leader's longitudinal position. The collision-avoidance constraint is  $x_{\text{ego}}(t) \leq x_{\text{lead}}(t) - \Delta$  for some positive stand-off  $\Delta$ ; the analysis is insensitive to  $\Delta \geq 0$ , so take  $\Delta = 0$  for the bound. The leader's clearance time is  $t_{\text{lead}} = (d_l + L_i)/v_l$ , which is exactly  $t_c^{\text{eff}}$  of Definition 21 (the small constant  $\delta$  prevents division by zero when  $v_l = 0$  and makes the inequality strict in the border case). If  $t_c^{\text{eff}} \geq t_y - \epsilon$ , the leader occupies the box at  $t = t_y - \epsilon$ , i.e.,  $x_{\text{lead}}(t_y - \epsilon) < d_l + L_i$ . By the collision-avoidance constraint,  $x_{\text{ego}}(t_y - \epsilon) \leq x_{\text{lead}}(t_y - \epsilon) < d_l + L_i \leq d_s + L_i$  (using  $d_l \leq d_s$ ), so the ego has not cleared the far boundary by the margin instant; moreover the ego has already crossed the stop line, since it cannot stop (see below). Hence the ego is blocked (Definition 3).

It remains to show the ego cannot abort. At the decision instant the gap to the leader is  $d_l < d_{\text{req}}(v_e)$ . By Theorem 12, any braking trajectory that halts the ego does so at a position at least  $d_{\text{req}}(v_e)$  beyond the decision point, i.e., at or beyond the leader's current position; since the leader is moving forward (or is stopped), the ego cannot come to rest strictly behind the leader without violating the collision-avoidance constraint. The ego therefore cannot stop in time and must proceed, which we have shown leads to blockage.  $\square$

### G. Proof of Theorem 25 (Cut-in warning)

*Proof.* By Definition 24, the intruding vehicle crosses the lane boundary at  $t = t_i = W_l/|v_{\text{lat}}|$ . If  $t_i < t_y$ , the intrusion occurs strictly before the signal transition. At the intrusion instant, the longitudinal separation between ego and intruder is at most  $d_a$  (the intruder is ahead at the ego's lane entry point; if it enters behind the ego the situation is not safety-relevant, so consider the ahead case). Since  $d_a < d_{\text{req}}(v_e)$ , by Theorem 12 the ego cannot brake to a halt before reaching the intruder's position. Two cases arise:

*Case 1 (ego brakes):* the ego cannot stop before the intruder, so it either collides with the intruder or, if the intruder is beyond the stop line, stops beyond the stop line and occupies the box. Both outcomes violate Definitions 3 or safe operation.

*Case 2 (ego proceeds):* the intruder now occupies the ego's lane and acts as a slow lead vehicle; by Theorem 22 with the intruder as leader (its gap  $d_a < d_{\text{req}}(v_e)$  and its speed implying a clearance failure), blockage follows.

In either case the kinematic constraints are violated and a warning is required.  $\square$

### H. Soundness of the Decision Pipeline

**Lemma 34** (Pipeline soundness). *Let  $\ell^*$  be the output of evaluate\_safety on any scene. Then  $\ell^* = \max_{\text{severity}} \{\ell_r\}$  over all rules  $r$  that fire, where the severity order is SAFE < CAUTION < WARNING < CRITICAL; if no rule fires,  $\ell^* = \text{SAFE}$ .*

*Proof.* The pipeline evaluates the rules in the fixed order [rule\_light, rule\_dilemma, rule\_lead, rule\_cutin, rule\_stale], and find\_map returns the first element that evaluates to Some. Inspection of each rule shows its maximum severity level: rule\_light and rule\_dilemma can emit CRITICAL (the highest); rule\_lead can emit CRITICAL (sub-rule 3a) or WARNING; rule\_cutin emits at most WARNING; rule\_stale emits at most CAUTION. Because the pipeline order is non-increasing in the severity bound of each rule, the first firing rule carries the maximum severity among all firing rules. If none fires, unwrap\_or(Safe) returns SAFE.  $\square$

**Corollary 35** (Red-light dominance). *If LightState::Red, the output is CRITICAL regardless of all other inputs.*

*Proof.* rule\_light maps Red to CRITICAL and is evaluated first; by Lemma 34 the output is the maximum severity, i.e., CRITICAL.  $\square$

## REFERENCES

- [1] D. Gazis, R. Herman, and A. Maradudin, “The problem of the amber signal light in traffic flow,” *Oper. Res.*, vol. 8, no. 1, pp. 112–132, 1960. <https://doi.org/10.1287/opre.8.1.112>
- [2] D. A. Redelmeier and R. J. Tibshirani, “Association between cellular-telephone calls and motor vehicle collisions,” *N. Engl. J. Med.*, vol. 336, no. 7, pp. 453–458, 1997. <https://doi.org/10.1056/NEJM199702133360701>
- [3] S. Shalev-Shwartz, S. Shammah, and A. Shashua, “On a formal model of safe and scalable self-driving cars,” *arXiv preprint arXiv:1708.06374*, 2017. <https://arxiv.org/abs/1708.06374>
- [4] J. Redmon, S. Divvala, R. Girshick, and A. Farhadi, “You only look once: Unified, real-time object detection,” in *Proc. IEEE CVPR*, 2016, pp. 779–788. <https://arxiv.org/abs/1506.02640>
- [5] N. Wojke, A. Bewley, and D. Paulus, “Simple online and realtime tracking with a deep association metric,” in *Proc. IEEE ICIP*, 2017, pp. 3645–3649. <https://arxiv.org/abs/1603.00831>
- [6] K. Behrendt and L. Novak, “A deep learning approach to traffic lights: Detection, tracking, and classification,” in *Proc. IEEE ICRA*, 2017, pp. 1244–1249. <https://doi.org/10.1109/ICRA.2017.7989163>
- [7] M. Althoff and J. M. Dolan, “Online verification of automated road vehicles using reachability analysis,” *IEEE Trans. Robot.*, vol. 30, no. 4, pp. 903–918, 2014. <https://doi.org/10.1109/TRO.2014.2312453>
- [8] A. Dosovitskiy, G. Ros, F. Codevilla, A. Lopez, and V. Koltun, “CARLA: An open urban driving simulator,” in *Proc. Conf. on Robot Learning*, 2017, pp. 1–16. <https://arxiv.org/abs/1711.04238>
- [9] P. A. Lopez et al., “Microscopic traffic simulation using SUMO,” in *Proc. IEEE ITSC*, 2018, pp. 2575–2582. <https://arxiv.org/abs/1802.02215>
- [10] C. Godard, O. Mac Aodha, M. Firman, and G. J. Brostow, “Digging into self-supervised monocular depth estimation,” in *Proc. IEEE ICCV*, 2019, pp. 3828–3838. <https://arxiv.org/abs/1806.01260>
- [11] J. Philion and S. Fidler, “Lift, splat, shoot: Encoding images from arbitrary camera rigs by implicitly unprojecting to 3D,” in *Proc. ECCV*, 2020, pp. 194–210. <https://arxiv.org/abs/2004.02903>
- [12] N. Matsakis and F. Klock, “The Rust language,” *ACM SIGAda Ada Letters*, vol. 34, no. 3, pp. 103–104, 2014. <https://doi.org/10.1145/2692956.2663188>
- [13] S. Klabnik and C. Nichols, *The Rust Programming Language*, 2nd ed. San Francisco, CA, USA: No Starch Press, 2019. <https://doc.rust-lang.org/book/>
- [14] V. Vovk, A. Gammernan, and G. Shafer, *Algorithmic Learning in a Random World*, 2nd ed. Cham, Switzerland: Springer, 2022. <https://doi.org/10.1007/978-3-031-06649-8>
- [15] C. V. Zegeer and R. C. Deen, “Green-extension systems at high-speed intersections,” *ITE J.*, vol. 48, no. 11, pp. 19–24, 1978. <https://scholar.google.com/scholar?q=%22Green-extension+systems+at+high-speed+intersections%22>
- [16] ISO, *Road vehicles – Functional safety*, ISO 26262:2018, 2018. <https://www.iso.org/standard/68383.html>

**Algorithm 3** Decision rules (src/rules.rs). Each returns a warning level or NIL.

```

1: function RULERED( $\ell$ )
2:   if  $\ell = \text{RED}$  then return CRITICAL
3:   end if
4:   return NIL
5: end function
6:
7: function RULEDILEMMA( $v_e, d_s, t_y$ )
8:    $d_{\text{req}} \leftarrow \text{STOPPINGDISTANCE}(v_e)$ 
9:    $t_c \leftarrow \text{CLEARANCETIME}(d_s, v_e)$ 
10:   $\text{no\_stop} \leftarrow (d_s \leq d_{\text{req}})$ 
11:   $\text{no\_clear} \leftarrow (t_c \geq t_y - \epsilon)$ 
12:  if  $\text{no\_stop} \wedge \text{no\_clear}$  then return CRITICAL
13:  end if
14:  return NIL
15: end function
16:
17: function RULELEAD( $v_e, d_l, v_l, \text{in\_box}, t_y$ )
18:   $d_{\text{req}} \leftarrow \text{STOPPINGDISTANCE}(v_e)$ 
19:  if  $v_l < 1.0 \wedge d_l < d_{\text{req}} \wedge \text{in\_box}$  then
20:    return CRITICAL  $\triangleright$  3a: leader stopped inside box
21:  end if
22:   $t_{\text{eff}} \leftarrow \text{CLEARANCETIME}(d_l, v_l)$ 
23:  if  $t_{\text{eff}} \geq t_y - \epsilon \wedge d_l < d_{\text{req}}$  then
24:    return WARNING  $\triangleright$  3b: clearance failure behind
    leader
25:  end if
26:  return NIL
27: end function
28:
29: function RULECUTIN( $D, v_e, t_y$ )
30:   $d_{\text{req}} \leftarrow \text{STOPPINGDISTANCE}(v_e)$ 
31:  for all  $o \in D$  where  $o.\text{class} \in \{2, 3, 5, 7\}$  do
32:    if  $o.\text{signal} \wedge o.\text{lane} \neq \text{same}$  then
33:      if  $o.\text{age} \geq 3 \wedge |o.v_{\text{lat}}| \leq 4.0$  then
34:         $t_i \leftarrow \text{INTRUSIONTIME}(o.v_{\text{lat}})$ 
35:        if  $t_i < t_y \wedge o.d < d_{\text{req}}$  then return
        WARNING
36:      end if
37:    end if
38:  end for
39:  return NIL
40: end function
41:
42:
43: function RULEYELLOW( $\ell, t_y$ )
44:  if  $\ell = \text{YELLOW} \wedge t_y < 2.5$  then return CAUTION
45:  end if
46:  return NIL
47: end function
48:
49: function RULESTALE( $\ell, v_e, d_s$ )
50:  if  $\ell = \text{GREEN} \wedge d_s > \text{STOPPINGDISTANCE}(v_e)$  then
51:    return CAUTION  $\triangleright$  Worst-case: green may end at
    any frame
52:  end if
53:  return NIL
54: end function

```

---

**Algorithm 4** Decision pipeline (src/decision.rs). Composes rules in descending severity.

---

```
1: function EVALUATESAFETY( $v_e, d_s, D, \ell, t_y$ )
2:    $l \leftarrow$  first  $o \in D$  with  $o.lane = \text{same}$ ,  $o.class \in \{2, 3, 5, 7\}$ 
3:   if  $l \neq \text{NIL}$  then
4:      $d_l \leftarrow l.d$ ,  $v_l \leftarrow l.v$ ,  $in\_box \leftarrow (l.d < L_i)$ 
5:   end if
6:
7:    $R \leftarrow [ \text{RULERED}(\ell),$ 
8:      $\text{RULEDILEMMA}(v_e, d_s, t_y),$ 
9:      $\text{RULELEAD}(v_e, d_l, v_l, in\_box, t_y),$ 
10:     $\text{RULECUTIN}(D, v_e, t_y),$ 
11:     $\text{RULEYELLOW}(\ell, t_y),$ 
12:     $\text{RULESTALE}(\ell, v_e, d_s) ]$ 
13:   for all  $r \in R$  do  $\triangleright$  Evaluated in severity order; first match wins
14:      $w \leftarrow r()$ 
15:     if  $w \neq \text{NIL}$  then return  $w$ 
16:   end if
17: end for
18: return SAFE
19: end function
```

---

---

**Algorithm 5** Binary entry point (src/main.rs).

---

```
1: Subcommand RUN( $source$ ):
2:    $F \leftarrow \text{CAMERACAPTURE}(source)$   $\triangleright$  CSI camera or video file
3:
4:   for all  $f \in F$  do
5:      $D \leftarrow \text{YOLODETECT}(f)$   $\triangleright$  ONNX Runtime, INT8 quantised
6:      $T \leftarrow \text{DEEPSORT}(D)$   $\triangleright$  Kalman filter + Hungarian association
7:      $w \leftarrow \text{EVALUATESAFETY}(v_e, d_s, T, \ell, t_y)$ 
8:
9:     if  $w \neq \text{SAFE}$  then DISPATCHALERT( $w$ )
10:    end if
11:
12:   end for
13:
14: Subcommand COLLECT( $source, rate$ ):
15:   Save frames at  $rate$  Hz as JPEGs for YOLO training
16:
17: Subcommand TRAIN:
18:   Prepare dataset  $\rightarrow$  GPU training  $\rightarrow$  ONNX export  $\rightarrow$  validate
```

---

---

**Algorithm 6** Project structure (Cargo.toml). Zero external dependencies for core logic.

---

```
1: Package: driving-civicsense-vision-model v0.1.0
2: Rust edition: 2024, MSRV 1.85+
3: License: GNU AGPL v3
4:
5: Core engine (zero dependencies):
6:   algebra, models, rules, decision
7:
8: Binary (runtime dependencies):
9:   onnxruntime – YOLO inference
10:  opencv – video capture and visualization
11:  tonic, prost – gRPC alert dispatch
12:  clap – CLI argument parsing
13:  serde, serde_yaml – configuration
```

---